


4. GLOBAL STATE AND SNAPSHOT RECORDING ALGORITHMS

Ex



Let acct. A recorded at t_0 : \$600

Can recorded at t_2 : \$50


C_{A2} recorded at t_2 : \$80

acct. B recorded at t_2 : \$120

\$850!!

{inconsistent}
A's state was
recorded before
\$50 was sent

error \$50



Global state: $\{ \bigcup_i PS_i, \bigcup_{i,j} CS_{ij} \}$
collection of local states of processes & channels

it is consistent iff:

C1: $send(m_{ij}) \in PS_i \wedge m_{ij} \in CS_{ij} \wedge recv(m_{ij}) \in PS_j$ (law of conservation of msgts.)

C2: $send(m_{ij}) \notin PS_i \wedge m_{ij} \notin CS_{ij} \wedge recv(m_{ij}) \notin PS_j$ (every effect has a cause)

consistent snapshot \rightarrow all recorded states of processes are concurrent.

Issues to be addressed in recording consistent snapshot:

I1: which messages to be recorded in snapshot

I2: when should a process take its snapshot


SNAPSHOTS IN CAUSAL CHANNELS

Causal ordering provides a builtin synchronization to messages. Snapshot algorithms for such systems are considerably simplified. (e.g. - do not send control msgs on every channel)

CO: $\text{Send}(m_{ij}) \rightarrow \text{Send}(m_{kj}) \rightarrow \text{recv}(m_{ij}) \rightarrow \text{recv}(m_{kj})$

• Process state recording (common) • channel state recording (Acharya-Badrinath, Agha-Venkatraman)

PROCESS STATE RECORDING



① initiator broadcasts token (inc. idseq)

② each process on $\text{recv}(\text{token}_i)$ captures local state PS_i and sends it to the initiator


③ terminated when initiator receives the snapshot recorded by each process

$\text{Send}(\text{token}) \rightarrow \text{recv}(\text{token}_i) \rightarrow \text{Send}(m_{ij})$

$\rightarrow \text{Send}(\text{token}) \rightarrow \text{recv}(\text{token}_i) \rightarrow \text{recv}(m_{ij})$

∴ $\text{Send}(m_{ij}) \notin PS_i \rightarrow \text{recv}(m_{ij}) \notin PS_j$ (C2)

ACHARYA-BADRINATH ALGORITHM



Each process P_i maintains:

- SENT_i = [msgs sent by P_i to others]

- RECD_i = [msgs received by P_i from others]

• $\text{RECD}_1 = [0, 0, 0]$ before $\text{recv}(\text{token}_1)$

$\text{SENT}_3 = [0, 1, 0]$ before $\text{recv}(\text{token}_3)$


$c_{31} = m_{31}$ ($\text{RECD}_3[3] + 1$ $\text{SENT}_3[1]$)

- When a process P_i records its local snapshot PS_i on $recv(token_i)$, it includes $SENT_i$ and $RECD_i$ in its local state before sending snapshot to the initiator.
- When algorithm terminates, initiator determines the state of channels in the global snapshot being assembled as follows:
 - $C_{1i} = \emptyset$ (initiator to others empty)
 - $C_{ij} = \{RECD_j[i] + 1, \dots, SENT_i[j]\}$ (msgs which have been sent, but not yet received)

Let $k = seq. no. (msg)$

- case $SENT_i[j] < k$: message was sent after $recv(token_i)$
 $send(msg) \notin PS_i \Rightarrow recv(msg) \notin PS_i$ (due to CO) (C2)
- case $RECD_j[i] < k \leq SENT_i[j]$: msg sent before $recv(token_i)$, but after $recv(token_j)$
 $msg \in C_{ij}$ as per rule (2) (C1)
- case $k \leq RECD_j[i]$: msg sent before $recv(token_i)$ and also before $recv(token_j)$
 $recv(msg) \in PS_j$ as $token_j$ was received after it

ALAGAR - VENKATESAN ALGORITHM



A message is referred to as:


- old: if $send(msg)$ precedes $send(token)$
 - new: otherwise (can be found using vector timestamps)
- old: $[m_1, m_2]$, new: $[m_3]$

(x_1, x_2, x_3 : terminate message)


NOTE: a process receives all old msgs before terminate msg due to causal ordering

- When a process P_i receives $token_i$, it takes its snapshot, initializes all its channels to empty, and returns "done" message to the initiator. Any further "old" msgs received are added to channel state.
 - after initiator receives "done" message from all processes, it broadcasts a "terminate" message.
 - a process stops the snapshot algorithm after receiving the "terminate" message.
- m_3 : new not added to channels by (1) (C2) m_1 : old after done added to channel m_2 : included in PS_2 , (C1)

SNAPSHOTS IN FIFO CHANNELS : CHANDY - LAMPORT ALGORITHM



this diagram does not appear to follow consistent snapshot with the algorithm ↑. is this a FIFO channel?



an initiator first executes marker sending rule. the algorithm terminates after each has recieved a marker on all of its incoming channels. recorded local snapshots can be put together to create the global snapshot in several ways.

- marker sending rule for P_i :
- ① process P_i records its state
- ② send marker on all outgoing channel, if not sent already.

- each process can send its local snapshot to the initiator.
- each process can spread the information to all outgoing channels (all can determine G_1)


- marker receiving rule for P_j :
- ① if P_j has not recorded its state:
 - (a) mark channel C as empty
 - (b) execute - "marker sending rule"
- ② else
 - record state of C as set of msgs recieved along C after recording state and before receiving marker along C .

- msg sent after marker not recorded (G_2) in channel state (FIFO channel) (C_2)
- msg sent before marker on channel (m_1) → after recording state → inc. in channel (C_1)
- (m_1) → before recording state → inc. in snapshot

SNAPSHOTS IN NON-FIFO CHANNELS

hellary algorithm uses msg inhibition to avoid inconsistency in global snapshot. when a proc. receives a marker, it immediately returns an acknowledgement. after P_i has sent marker to P_j it does not send it any msg until it receives an ack. or a marker from P_j .

LAI - YANG ALGORITHM



① every process is initially white, execute "marker sending" rule when it turns red

② white processes send white msgs, and red ones send red msgs

③ every white process takes a snapshot when it receives a red msg. $\langle c_2 \rangle$

$\hookrightarrow c_2$ send $(m_{red}) \notin PS_i$ and recv $(m_{red}) \notin P_{fbi}$

c_1 send $(m_{white}) \in P_{fbi} \rightarrow$ recv $(m_{white}) \in PS_i$ ④ $m_{white} \in G_{ij}$