

DISTRIBUTED MUTUAL EXCLUSION ALGORITHMS

mutual exclusion ensures that concurrent access of processes to a shared resource (or data) is serialized.

• token-based

critical section
↓

process holding unique token can enter CS.

algs differ by how they search for token.

• non-token based

2 or more rounds of messages are exchanged

i. to determine who can enter CS next, a process

e enters CS when an assertion becomes true

ii (and false for everyone else).

• quorum based

i. subsets of processes (quorums) are formed

c in such a way that when 2 processes

v request CS, at least one process receives

both requests and is responsible to make

sure only one process gets CS.

algorithm performance in load:

low load \rightarrow seldom more than one request for CS.

⑤ heavy load \rightarrow always a pending request for CS.

(a process is seldom in idle state in heavy load)

algorithm best and worst case performance:

most algs \rightarrow best response time = $2T + E$ round trip msg delay

often, best and worst cases coincide with low and high loads.

process P_i in one of 3 states:

- requesting CS (wait)
- executing CS
- neither (idle)


algorithm requirements:

- safety: only 1 process in CS
- liveness: no deadlock or starvation
- fairness: CS executed in order of arrival

performance metrics:

- message complexity
no of msgs per CS execution
- synchronization delay, so
time gap between CS exit & entry.
- response time
time gap b/w CS request & exit.
- system throughput

rate of CS execution $\frac{1}{SD + E}$



LAMPORT'S ALGORITHM

processes use scalar clocks and request-queue to perform mutual exclusion. request-queue is ordered by timestamps. process at the top of request-queue enters CS. requires FIFO channels for fairness, and mut. ex.

Proof: mutual exclusion achieved. (by contradiction)

assume P_1, P_2 in CS \rightarrow ①

assume P_1 REQUEST $<$ P_2 REQUEST (ts) \rightarrow ②

① \rightarrow P_2 sent REPLY to P_1

\rightarrow P_2 has P_1 REQUEST before P_1 REPLY (FIFO channel)

② \rightarrow P_1 REQUEST at top of P_2 request-queue

but ① \rightarrow P_2 REQUEST at top of P_2 request-queue.

contradiction \rightarrow mutual exclusion achieved.

proof: fairness achieved (by contradiction)

assume P_2 in CS before P_1 \rightarrow ①

assume P_1 REQUEST $<$ P_2 REQUEST (ts) \rightarrow ②

① \rightarrow P_2 received P_1 's REPLY

\rightarrow P_2 received P_1 REQUEST (FIFO channel)

② \rightarrow P_1 REQUEST at top of P_2 request-queue

but ① \rightarrow P_2 REQUEST at top of P_2 request-queue.

contradiction \rightarrow fairness achieved.

optimization:

if P_i , P_i REQUEST $<$ P_j REQUEST

\rightarrow no need to REPLY P_j

$$\rightarrow O(2(n-1)) \leftrightarrow O(2(n-1))$$

CS request:

- P_i broadcasts REQUEST(t_i, i). (inc. itself)
- P_j on receiving REQUEST(t_i, i) returns timestamped REPLY to P_i and adds REQUEST to request-queue.


execute CS:

L1: P_i received REPLY with larger timestamp from all.

L2: P_i 's REQUEST at top of request-queue.

release CS:

- P_i on CS exit broadcasts timestamped RELEASE to all (inc. itself).
- P_j on receiving RELEASE, removes P_i 's REQUEST from request-queue.



RICART - AGRAWALA ALGORITHM

Processes use scalar clocks and request deferred array to perform mutual exclusion without a FIFO channel.

Proof: mutual exclusion achieved (by contradiction)

assume P_1, P_2 in CS — ①

assume $P_1 \text{ REQUEST} < P_2 \text{ REQUEST}$ (i.e.) — ②

① $\Rightarrow P_2$ has P_1 REPLY

$\Rightarrow P_1$ not in CS & $P_1 \text{ REQUEST} > P_2 \text{ REQUEST}$

contradiction \Rightarrow mutual exclusion achieved.



msg complexity = $O(n(n-1))$

request CS:

- P_i broadcasts timestamped REQUEST.
- P_j on receiving REQUEST from P_i sends REPLY only if P_j is not executing CS and $P_j \text{ REQUEST}$ timestamp is larger than $P_i \text{ REQUEST}$. else, reply is deferred and P_j sets $RD_j[i] = 1$.

execute CS:

- P_i received REPLY from all.

release CS:

- P_i sends all deferred REPLY messages: $\forall RD_j[i] = 1$, P_i sends REPLY to P_j and sets $RD_j[i] = 0$.

MAEKAWA'S ALGORITHM

first quorum based algorithm. a process only

requests a subset of processes (quorum). two

simultaneously requesting processes withing their

quorum will have atleast 1 common process, which

REPLYs to only one of the two.

mutual exclusion achieved.

conditions for request sets

request set (quorum)

$$M1: R_i \cap R_j \neq \emptyset \quad \forall i \neq j$$

$$M2: P_i \in R_i \quad \forall i$$

$$M3: |R_i| = K \quad \forall i$$

$$M4: P_i \in K \text{ no. of } R_i\text{'s} \quad \forall i, j$$

$$N = K(K-1) + 1$$

no. of processes

$$|R_i| = \sqrt{N}$$

(theory of projective planes)

proof: mutual exclusion achieved (by induction) request CS:

assume P_i, P_j in CS — ①

assume $R_i \cap R_j = \{P_{ij}\}$ — ②

① $\rightarrow P_i$ got REPLY from all in R_i

② $\rightarrow P_j$ got REPLY from all in R_j

③ $\rightarrow P_{ij}$ sent REPLY to P_i and P_j .


contradiction \rightarrow mutual exclusion achieved.

msg complexity: $O(\sqrt{N})$ REQUEST
REPLY
RELEASE

maekawa's algorithm can deadlock because

a process is locked by other processes, and

requests are not prioritized by timestamps.



- in response to INQUIRE(j) from P_{ij} , P_i sends YIELD(i) to P_{ij} if P_i has received FAILED from a process in R_i , and if it sent YIELD to any one but not received new REPLY from it.

msg complexity: $O(\sqrt{N})$

request CS:

- P_i sends REQUEST(i) to all processes in R_i .
- P_j on receiving REQUEST(i), sends REPLY(j) only if it hasn't sent REPLY to a process since last RELEASE msg. else, it queues REQUEST(i) for later consideration.

execute CS:

- P_i enters CS after it received REPLY from every process in R_i .

release CS:

- P_i sends RELEASE(i) to every site in R_i .
- P_j on receiving RELEASE(i), sends REPLY to next process in its queue after deleting P_i . if queue is empty, it updates its state to no REPLY sent since last RELEASE.

handling deadlocks:

- when REQUEST(i, j) from P_i blocks at P_{ij} because P_{ij} sent REPLY to P_i , P_{ij} sends FAILED(j) to P_i if its request has lower priority, else P_{ij} sends INQUIRE(j) to P_i .
- in response to YIELD(i) from P_i , P_{ij} assumes as if it has been released by P_i , places request of P_i at appropriate location in queue, and sends a REPLY(j) to process at top of queue.